

Learning PostgreSQL as a SQL Server User

Grant Fritchey

Grant Fritchey

DevOps Advocate

Microsoft PostgreSQL MVP
AWS Community Builder

 @gfritchey

 Grant.Fritchey@Red-Gate.com

 /in/gfritchey

 scarydba.com



Agenda

- Why PostgreSQL?
- A new language
- What's the same?
- What's different?
- What's confusing?
- How to go about learning PostgreSQL
- Where to go for more



Why PostgreSQL?

Licensing

- Extremely permissive, open-source license
- Core features can never be paywalled
- Most extensions piggyback on the PostgreSQL License



Extensibility



- >30 hooks available to extensions
- Any supported language, although SQL, C, and Rust are most popular

Extensibility

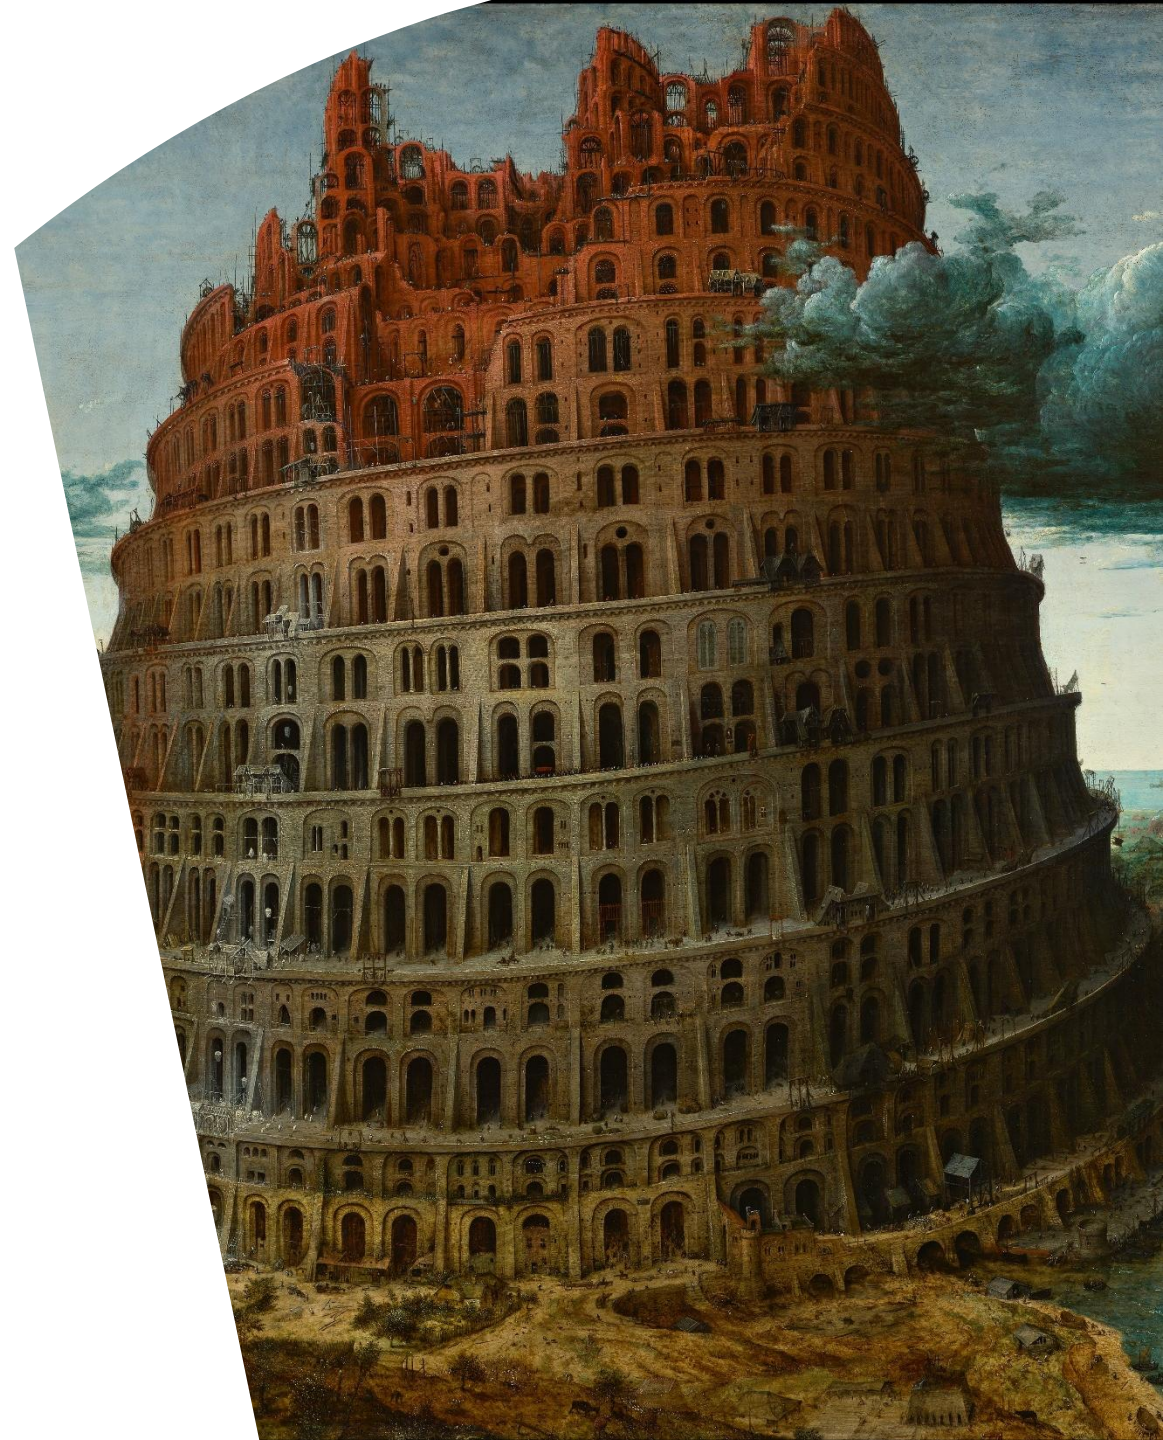


- Not just for extending functionality
- Package commonly used SQL functions/procedures
- Versioning required and helpful

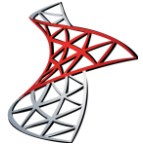
A New Language

Confusion

- Words do not have the same meaning
- Take the time to learn the new terms
- It makes learning everything else easier



Mapping Terms



SQL Server

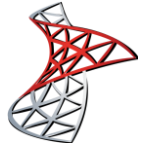


PostgreSQL

Instance

Cluster

Mapping Terms



SQL Server

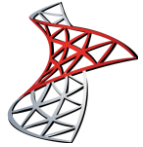
User



PostgreSQL

Role

Mapping Terms



SQL Server

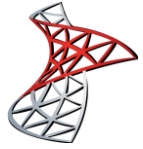


PostgreSQL

Transaction
Log

Write
Ahead Log
(WAL)

Mapping Terms



SQL Server

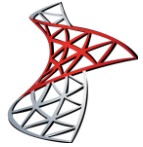


PostgreSQL

Row

Tuple

Mapping Terms



SQL Server

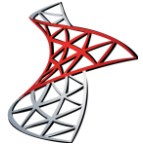


PostgreSQL

Procedure

Function

Mapping Terms



SQL Server



PostgreSQL

Function

Procedure

What's The Same?

Core Concepts



- Relational database management system (well, Object Relational)
- SQL
- Normalization

Common Across Both



- Database
- Schema
- Table
- Primary key
- Foreign key

Common Across Both



- Column
- View
- Index
- Trigger
- Constraint

Common Across Both



- Query
- SELECT
- INSERT...
- In fact, strict adherence to the ANSI SQL standard (mostly)

What's Different?

Case Rules

- All object names are lower case
- Unless you double “Quote” them
- `commonly_use_snake_case`



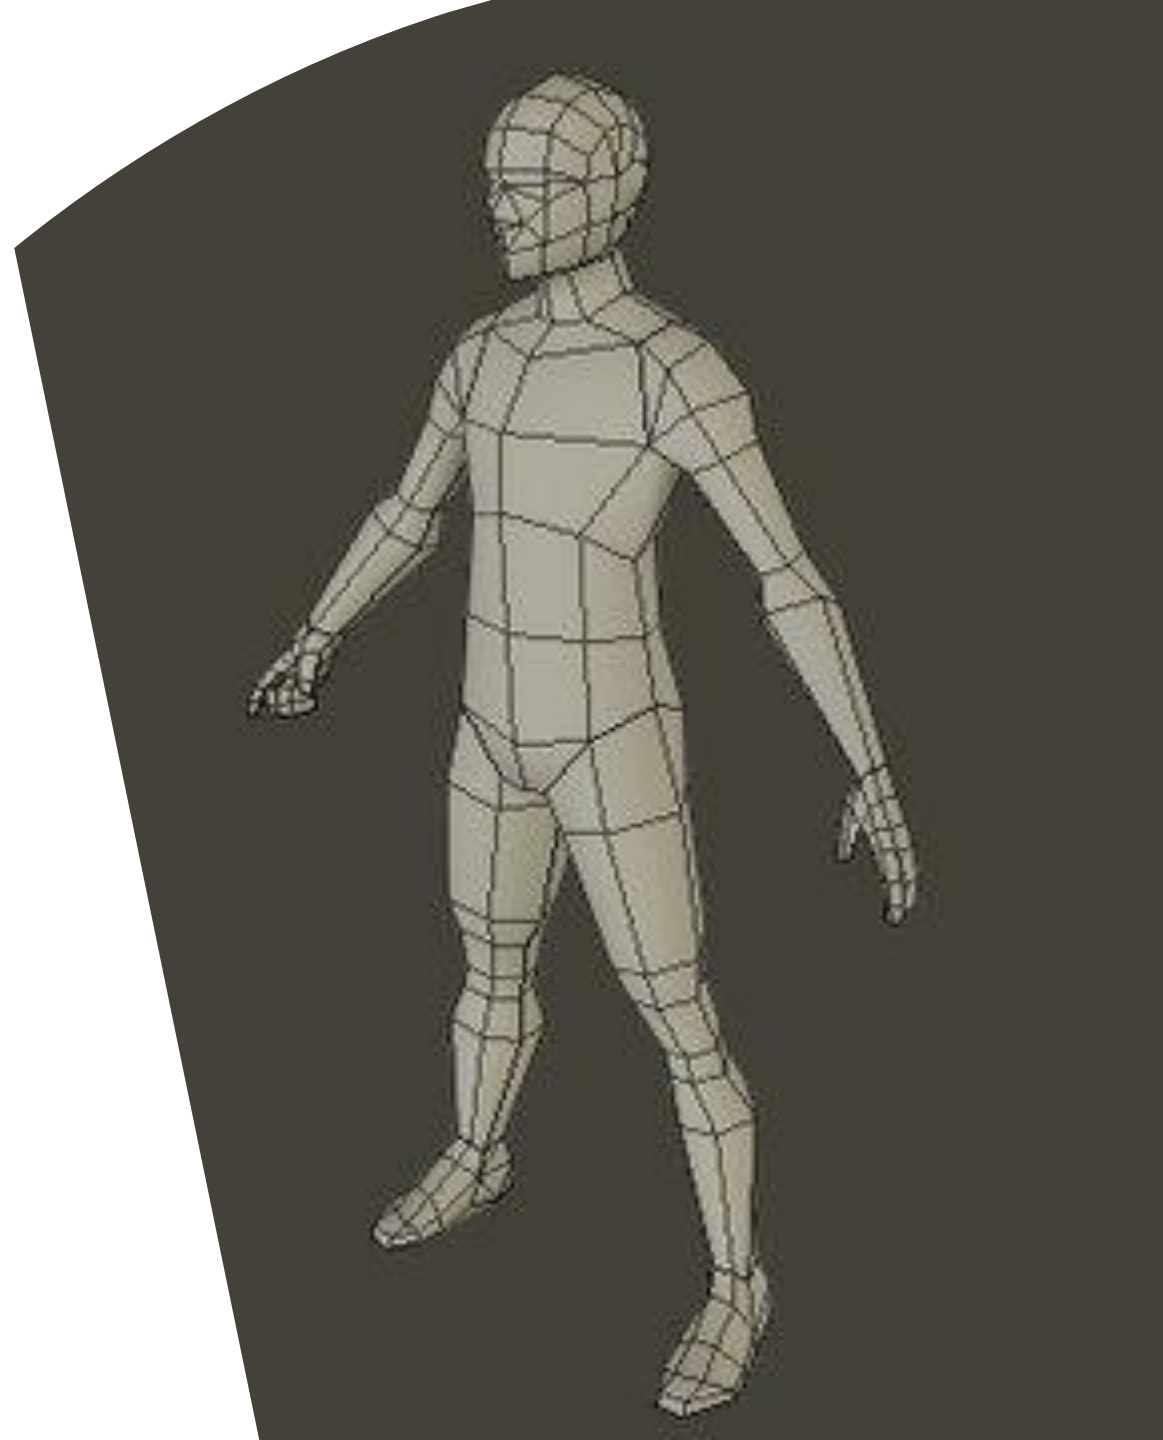
SQL

- LIMIT instead of TOP
- Ordinals instead of names
- DATE_PART instead of specific functions
- LATERAL instead of CROSS
APPLY



Database Templates

- Think, model db
- template0 – leave this alone
- template1 – customizable
- Custom templates
- You can choose the template



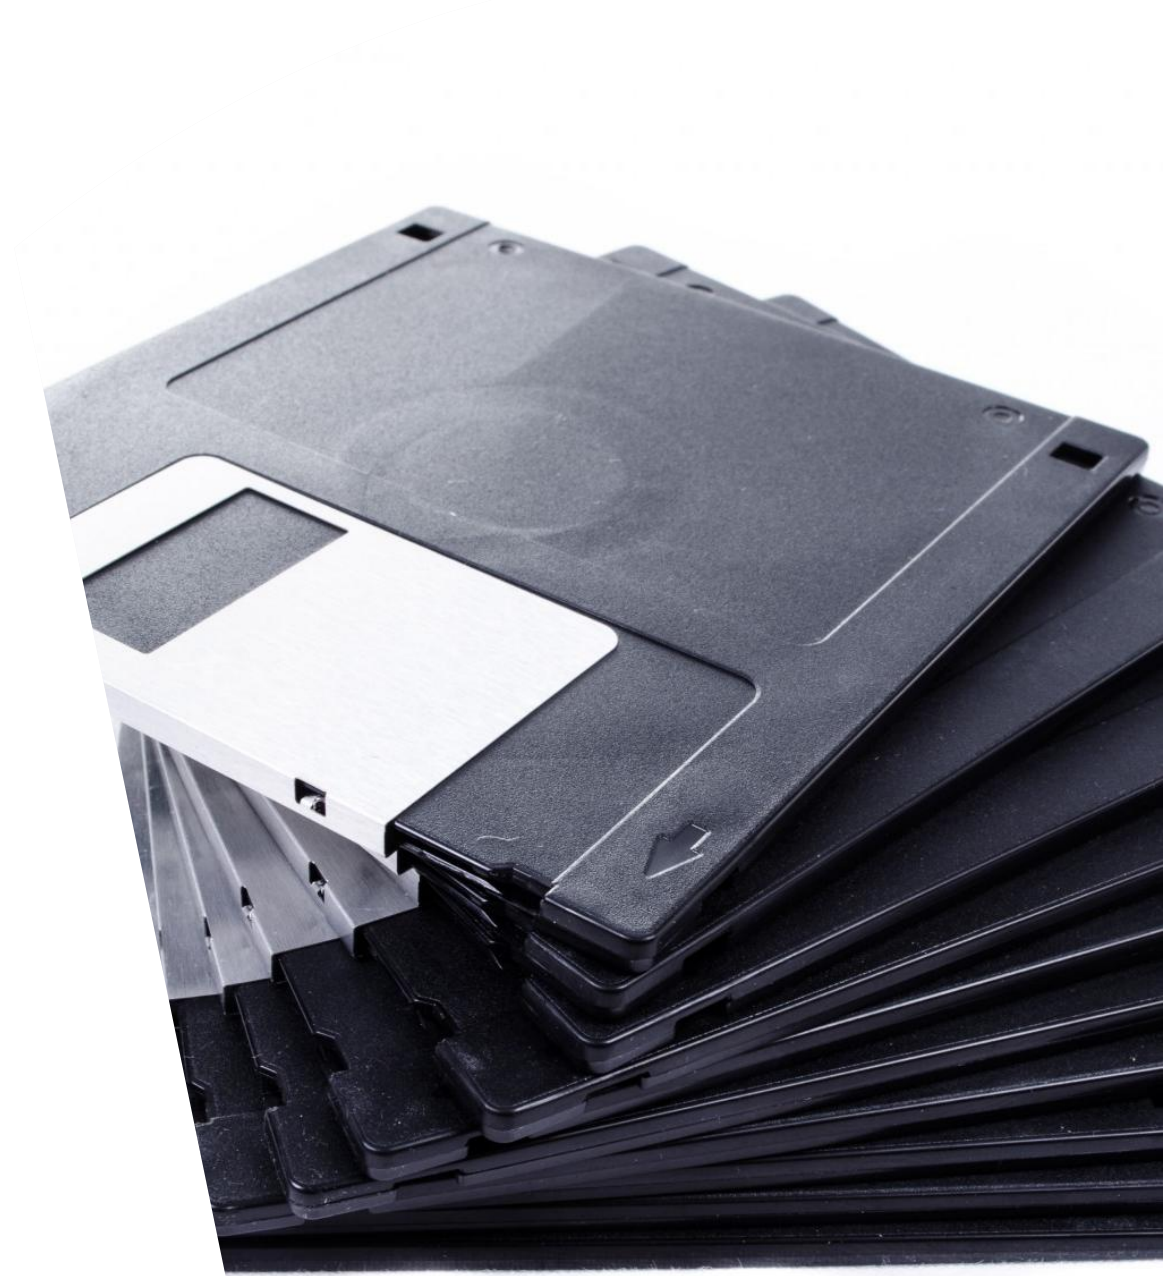
CREATE DATABASE

- Defaults to template1
- You can pick the template
- You can define the owner



Backups & Restores

- `pg_dump`, `pg_dumpall`
- `pg_basebackup`
 - `-- incremental` (version 17+)
- WAL archiving for point-in-time recovery (more later)
- `pg_restore`
- Extensions



Data Types

- Twenty (20) Categories of data types
- Depending on how you count, ~125 data types



Data Type Behaviors

- Fixed or variable length, depending on type
- Most can be NULL
- DEFAULT values can be defined
- Can be treated as objects, even nested



Indexes

- PostgreSQL tables are heaps
- No concept of a clustered index
- CREATE and DROP CONCURRENTLY
- Must be table owner to add indexes



Index Types

- B-Tree
- Hash
- GiST
- SP-GiST
- GIN
- BRIN



Core Procedural PostgreSQL Languages



- SQL
- PL/pgSQL
- PL/Python
- PL/TcL
- PL/Perl

Server Metrics



- Error logs
- Cumulative Statistics System
 - Off by default
 - Reset on crash, failover, or point in time restore
 - Can be manually reset
 - Aggregations only

Cumulative Statistics System



- `pg_stat_activity` – current activity on server
- `pg_stat_database` – activity & size of database
- `pg_stat_*_tables` – activity & size of table
 - `all`
 - `sys`
 - `user`
- `pg_stat_*_indexes` – guess
- `pg_statio_*_tables` – specific info on blocks & storage

Query Metrics



- EXPLAIN ANALYZE
- pg_stat_statements
 - Disabled by default

What's Confusing?

Restore to Point In Time



1. Run `pg_switch_wal`
2. Stop the server
3. Clean data directory or make a new one
4. Restore the Base Backup files
5. Remove all files from `pg_wal/`
6. Edit `postgresql.conf` for restore command and point in time
7. Create a file called `recovery.signal` in the cluster data directory
8. Start the service

Functions



- Implemented for returning data from PostgreSQL database
- Similar to Oracle functions or SQL Server procedures

Using Function Example

SELECT

g1.genre_id,

g1.genrename

FROM

genre_list() **AS** *g1*;

Procedures



- Procedures are meant to perform actions, as opposed to functions, which return data
- Similar to Oracle procedures, or SQL Server functions

Procedure Example

```
CREATE OR REPLACE PROCEDURE new_genre (genrename TEXT)
AS $$
BEGIN
    INSERT INTO bluebox.film_genre
        (name) VALUES (genrename);
END
$$
LANGUAGE plpgsql;
```


Vacuum



- Cleans dead tuples
- Reduces bloat
- Updates statistics
- Guaranteed, you need to run it more frequently

How To Learn PostgreSQL

Supported Operating Systems



- Linux
- Windows
- Containers
- Cloud

Packaged Bits

- Core
 - For the “hackers”
- Packages
 - OS
 - Special functions



Containers

- Easiest way to get started
- Get in the habit of setting up volumes
- Multiple forks available



Cloud

- Azure
- Google Cloud
- AWS
- EnterpriseDB
- Heroku



PostgreSQL Tools

psql

- Always available (almost)
- Command line
- Can work remote

```
hell 7.4.5
C:\Users\grant> psql -?
psql is the PostgreSQL interactive terminal.

Usage:
    psql [OPTION]... [DBNAME [USERNAME]]

General options:
    -c, --command=COMMAND      run only single command
    -d, --dbname=DBNAME        database name to connect to
    -f, --file=FILENAME        execute commands from file
    -l, --list                  list available databases
    -v, --set=, --variable=NAME=VALUE set psql variable NAME to VALUE
                                (e.g., -v ON_ERROR_STOP=1)
    -V, --version              output version information
    --no-psqlrc                do not read startup file
    --single-transaction        execute as a single transaction
    --help[=options]           show this help, then exit
    --help=commands            list backslash commands
    --help=variables            list special variables, etc.

Input and output options:
    -e, --echo-all            echo all input from script
    -E, --echo-errors          echo failed commands
    -q, --echo-queries          echo commands sent to server
    -Q, --echo-hidden          display queries that internal commands
                                execute
    -o, --log-file=FILENAME    send session log to file
    -r, --readline              disable enhanced command processing
```

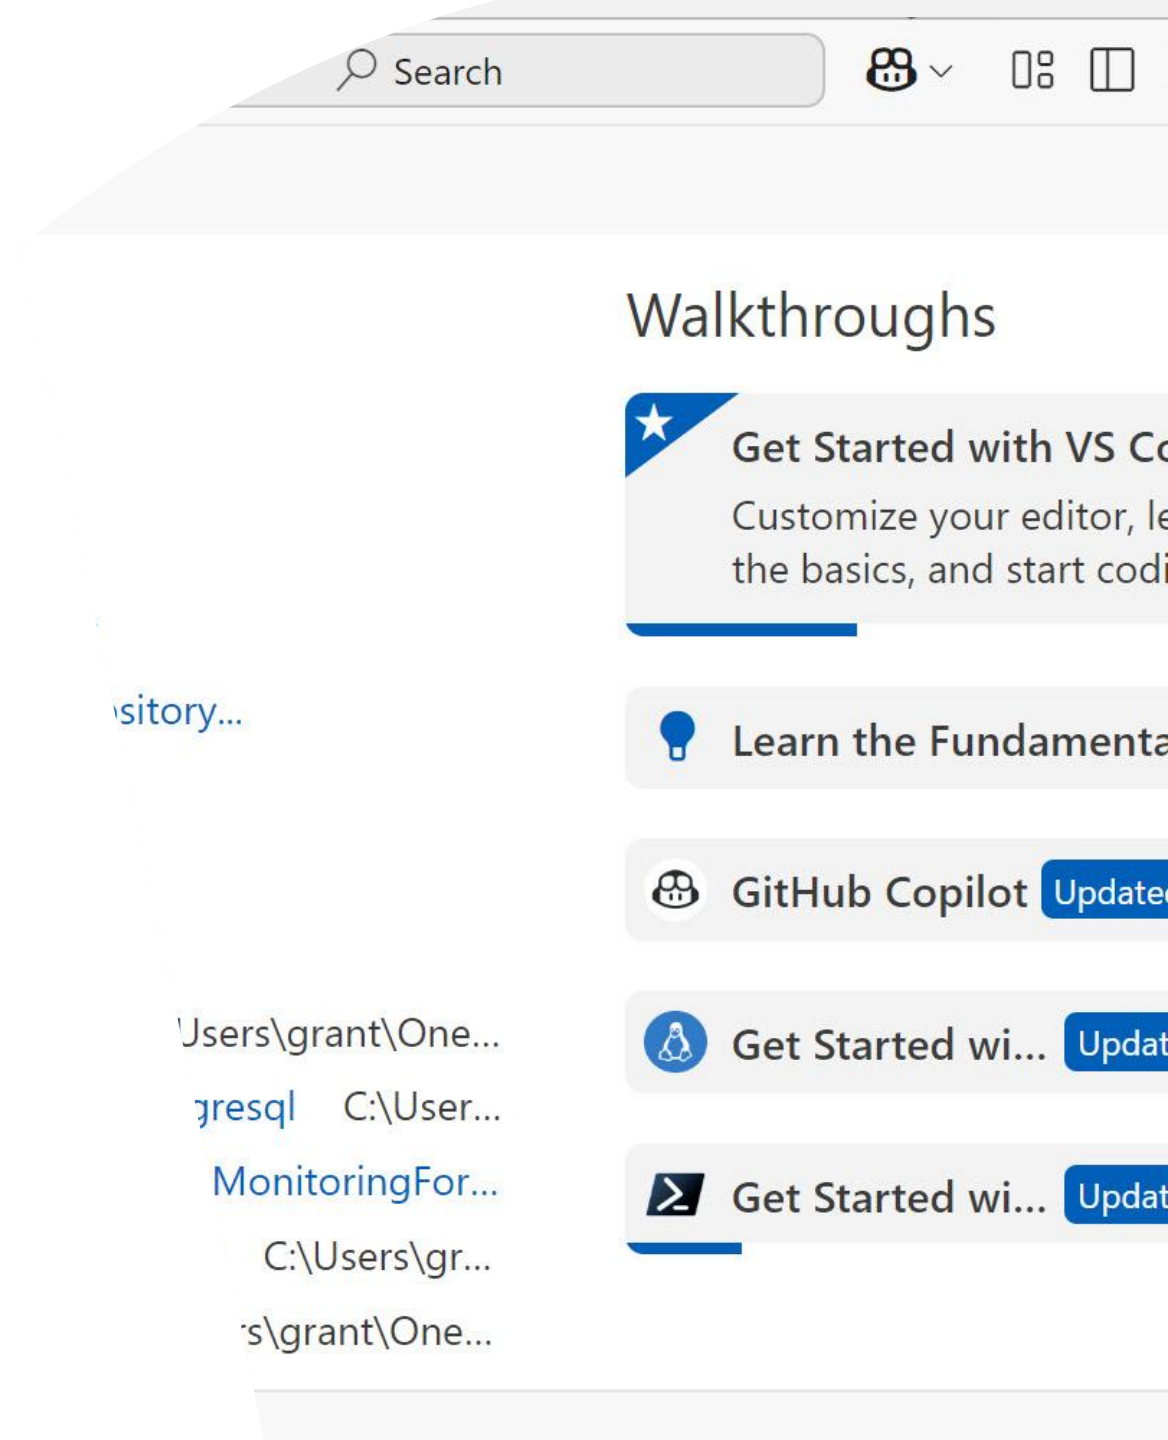
pgAdmin

- Unofficial, "Official" PostgreSQL GUI
- Not installed by default
- Web version for remote work
- Extensions



VS Code

- PostgreSQL Plugins
- Good for basics
- Good integration with source control



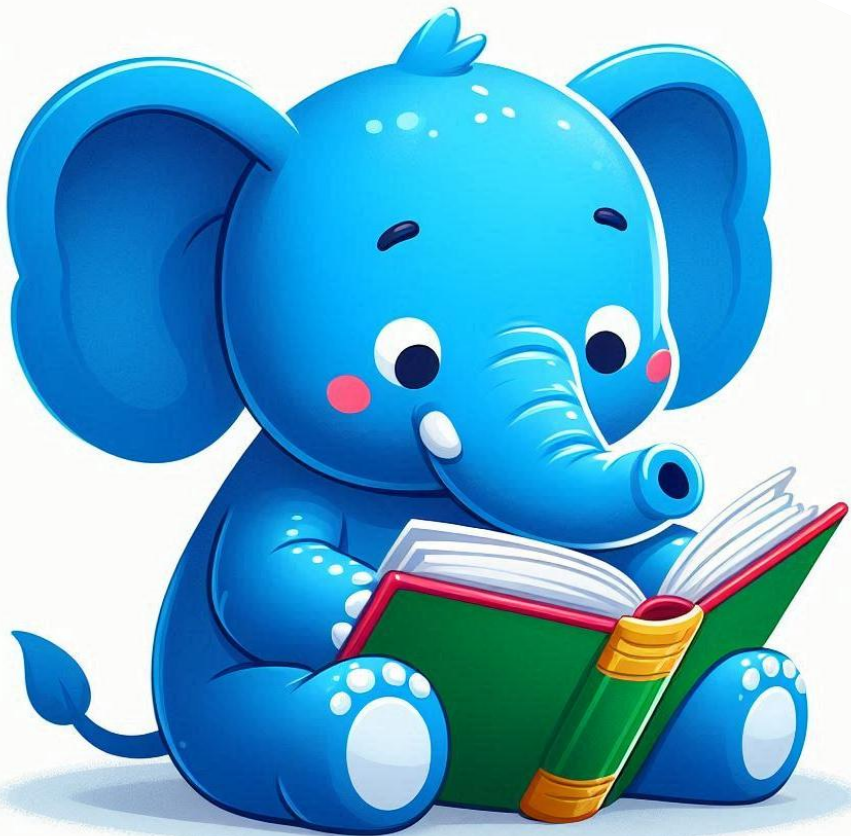
DBeaver

- Open source
 - Also paid
- Multi-platform
 - Good and bad
- Extensions
- Some code completion



More Learning

PostgreSQL Documentation



- <https://www.postgresql.org/docs/>
- Pay attention to this:
 - [https://www.postgresql.org/docs/
current/dml-delete.html](https://www.postgresql.org/docs/current/dml-delete.html)

Books



- [Introduction to PostgreSQL](#) – Ryan Booz, Grant Fritchey
- [Art of PostgreSQL](#) – Dimitri Fontaine
- [Database Administration](#) – Craig Mullins
- [PostgreSQL Query Optimization](#) - Henrietta Dombrovskaya, Boris Novikov and Anna Bailliekova

Other Events



- [PostgreSQL User Groups \(PUG\)](#)
- [MeetUp](#)
- [PGDay](#)
- pgConf (same link as PGDay)

Online



- [Planet PostgreSQL](#)
- [PostgreSQL Slack](#)
- Cooper Press [Email List](#)
- #pghelp on X
- Usual suspects

Feedback!

Please!



Grant Fritchey

DevOps Advocate

Microsoft PostgreSQL MVP
AWS Community Builder

 @gfritchey

 Grant.Fritchey@Red-Gate.com

 /in/gfritchey

 scarydba.com

